

```
!pip install requests beautifulsoup4 pandas
```

```

Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (2.32.3)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.11/dist-packages (4.13.4)
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (2.2.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests) (2.2.3)
Requirement already satisfied: certifi<2025.12.31, >=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests) (2025.11.12)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.11/dist-packages (from beautifulsoup4) (2.5)
Requirement already satisfied: typing-extensions>=4.0.0 in /usr/local/lib/python3.11/dist-packages (from beautifulsoup4) (4.13.2)
Requirement already satisfied: numpy>=1.23.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.9.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2-) (1.17.0)

```

```
!pip install yfinance
```

✓ 1. Data Preparation

- Web scraping exchange rate data and save to a certain format.

```

import requests
from bs4 import BeautifulSoup
import pandas as pd
from datetime import datetime

```

```

class CurrencyMonitor:
    def __init__(self):
        self.fetch_exchange_rate()

# update the url for the corresponding currency rate page, this only scraping CAD and CNY
def fetch_exchange_rate(self):
    url = "https://g.co/finance/CAD-CNY"
    print("Fetching...")
    #headers = {}
    response = requests.get(url)
    if response.status_code == 200:
        soup = BeautifulSoup(response.text, 'html.parser')
        soup_format = soup.prettify()
        #holders = soup.find_all('div', class_='YmKec fxKbKc')
        div_element = soup.find('div', {'data-source': 'CAD', 'data-target': 'CNY'})
        print(div_element)
        if div_element:
            # print(div_element)
            last_price = div_element.get('data-last-price')
            last_market_timestamp = div_element.get('data-last-normal-market-timestamp')
            unix_timestamp = int(last_market_timestamp)
            last_timestamp_converted = datetime.fromtimestamp(unix_timestamp).strftime('%Y-%m-%d %H:%M:%S')
            print(f"The last price is: {last_price} for date: {last_timestamp_converted}")
            return(last_timestamp_converted, last_price)
        else:
            print("No matching element")
            # In case if you want to output the reponse to a separate file instead showing in terminal, uncomment below
            # print("Writing to the text file response.txt...")
            # f = open('response.txt', 'w')
            # f.write(soup_format)
            # f.close()

```

CurrencyMonitor()

```
➦ Fetching...
<div data-entity-type="3" data-last-normal-market-timestamp="1754626464" data-last-price="5.2274519999999995" data-last-price-time="1754626464">
  The last price is: 5.2274519999999995 for date: 2025-08-08 04:14:24
<__main__.CurrencyMonitor at 0x7b7cebde9950>
```

#google page only has real time data, need to go to yahoo finance for historical data
import yfinance as yf

```
# Get last 6 months of daily data
df = yf.download("CADCNX=X", start="2025-02-01", end="2025-08-01", interval='1d')
df = df[['Close']]
df.reset_index(inplace=True)
df.columns = ['Date', 'Rate']
```

```
df.to_csv("cad_cny_history.csv", index=False)
print(df.tail())
```

```
➦ /tmp/ipython-input-3523705728.py:5: FutureWarning: YF.download() has changed argument auto_adjust default to True
  df = yf.download("CADCNX=X", start="2025-02-01", end="2025-08-01", interval='1d')
[*****100%*****] 1 of 1 completed      Date      Rate
122 2025-07-25  5.243906
123 2025-07-28  5.220485
124 2025-07-29  5.226109
125 2025-07-30  5.209860
126 2025-07-31  5.190529
```

df.head()

➦

	Date	Rate
0	2025-02-03	4.883466
1	2025-02-04	4.989753
2	2025-02-05	5.010937
3	2025-02-06	5.078317
4	2025-02-07	5.093781

✓ 2. Model Forecasting

✓ A. Using prophet: An open-source time series forecasting tool developed by Meta. It is a Bayesian statistical modelin.

- https://facebook.github.io/prophet/docs/quick_start.html

```
!pip install prophet
```

```
from prophet import Prophet
```

```
# Load data
```

```
df = pd.read_csv('cad_cny_history.csv')
```

```
df = df[['Date', 'Rate']]
```

```
df.columns = ['ds', 'y'] # ds:datestamp, y:target value
```

```
# Model
```

```
model = Prophet()
```

```
model.fit(df)
```

```
# Forecast
```


```
future = model.make_future_dataframe(periods=14) #forecast 14 days
```

```
forecast = model.predict(future)
```


```
# Plot
```

```
fig = model.plot(forecast)
```

```
➞ INFO:prophet:Disabling yearly seasonality. Run prophet with yearly_seasonality=True to override this.
INFO:prophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to override this.
DEBUG:cmdstanpy:input tempfile: /tmp/tmpba5dfjsc/lbv2uetk.json
DEBUG:cmdstanpy:input tempfile: /tmp/tmpba5dfjsc/3u7upw6b.json
DEBUG:cmdstanpy:idx 0
DEBUG:cmdstanpy:running CmdStan, num_threads: None
DEBUG:cmdstanpy:CmdStan args: ['/usr/local/lib/python3.11/dist-packages/prophet/stan_model/prophet_model.bin', '
04:16:07 - cmdstanpy - INFO - Chain [1] start processing
INFO:cmdstanpy:Chain [1] start processing
04:16:07 - cmdstanpy - INFO - Chain [1] done processing
INFO:cmdstanpy:Chain [1] done processing
```




```
forecast_14 = forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']]
print(forecast_14)
```



	ds	yhat	yhat_lower	yhat_upper
0	2025-02-03	4.979065	4.952071	5.004572
1	2025-02-04	4.995211	4.970576	5.021063
2	2025-02-05	5.011180	4.984458	5.037397
3	2025-02-06	5.030143	5.002914	5.056876
4	2025-02-07	5.050540	5.023828	5.077102
...
136	2025-08-10	5.155130	5.116527	5.192476
137	2025-08-11	5.217662	5.176187	5.256011
138	2025-08-12	5.218215	5.174286	5.260028
139	2025-08-13	5.218591	5.176561	5.266522
140	2025-08-14	5.221961	5.169978	5.272417


[141 rows x 4 columns]

```
fig2 = model.plot_components(forecast)
```



```
# Interactive figure
from prophet.plot import plot_plotly, plot_components_plotly

plot_plotly(model, forecast)
```



```
start_rate = forecast['yhat'].iloc[-14-1]
end_rate = forecast['yhat'].iloc[-1]
rate_change = end_rate - start_rate
trend = "increase" if rate_change > 0 else "decrease" if rate_change < 0 else "remain stable"

start_day = forecast['ds'].iloc[-14-1]
end_day = forecast['ds'].iloc[-1]

summary = (
    f"The CAD to CNY exchange rate is forecasted to {trend} by "
    f"{abs(rate_change)/start_rate*100:.2f}% over the next 14 days, "
    f"from {start_rate:.4f} to {end_rate:.4f} for date: {start_day} to {end_day}."
)

print("Prediction Summary:")
print(summary)
```

Prediction Summary:

The CAD to CNY exchange rate is forecasted to decrease by 0.19% over the next 14 days, from 5.2318 to 5.2220 for

B. Using ARIMA

- see separate notebook

C. Using LSTM/RNN/Transformer

- see separate notebook

3. Using LLM to explain insights of generated prediction results

Install Hugging Face Transformers & Datasets

!pip install transformers accelerate bitsandbytes huggingface_hub

```

Requirement already satisfied: transformers in /usr/local/lib/python3.11/dist-packages (4.55.0)
Requirement already satisfied: accelerate in /usr/local/lib/python3.11/dist-packages (1.9.0)
Collecting bitsandbytes
  Downloading bitsandbytes-0.46.1-py3-none-manylinux_2_24_x86_64.whl.metadata (10 kB)
Requirement already satisfied: huggingface_hub in /usr/local/lib/python3.11/dist-packages (0.34.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from transformers) (3.18.0)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (2.0.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages (from transformers) (6.0.2)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (2024.11.6)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from transformers) (2.32.3)
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.20.1)
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.5.2)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.11/dist-packages (from transformers) (4.67.1)
Requirement already satisfied: psutil in /usr/local/lib/python3.11/dist-packages (from accelerate) (5.9.5)
Requirement already satisfied: torch>=2.0.0 in /usr/local/lib/python3.11/dist-packages (from accelerate) (2.6.0+cu124)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub) (2024.12.0)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub) (4.12.2)
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub) (1.1.7)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->accelerate) (3.4.2)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->accelerate) (3.1.3)
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch>=2.0.0->accelerate)
  Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch>=2.0.0->accelerate)
  Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch>=2.0.0->accelerate)
  Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch>=2.0.0->accelerate)
  Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch>=2.0.0->accelerate)
  Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch>=2.0.0->accelerate)
  Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-curand-cu12==10.3.5.147 (from torch>=2.0.0->accelerate)
  Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch>=2.0.0->accelerate)
  Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparselt-cu12==0.6.2 (from torch>=2.0.0->accelerate)
  Downloading nvidia_cusparselt_cu12-0.6.2-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->accelerate) (0.6.2)
Collecting nvidia-nccl-cu12==2.21.5 (from torch>=2.0.0->accelerate)
  Downloading nvidia_nccl_cu12-2.21.5-py3-none-manylinux2014_x86_64.whl.metadata (1.8 kB)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->accelerate) (12.4.127)
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch>=2.0.0->accelerate)
  Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->accelerate) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->accelerate) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1) (1.3.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests) (2025.8.3)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from jinja2->torch>=2) (3.0.2)
Downloading bitsandbytes-0.46.1-py3-none-manylinux_2_24_x86_64.whl (72.9 MB)
 72.9/72.9 MB 14.0 MB/s eta 0:00:00
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl (363.4 MB)
 363.4/363.4 MB 5.6 MB/s eta 0:00:00
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (13.8 MB)
 13.8/13.8 MB 122.4 MB/s eta 0:00:00
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (24.6 MB)
 24.6/24.6 MB 66.9 MB/s eta 0:00:00
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (883 kB)
 883.7/883.7 kB 53.9 MB/s eta 0:00:00
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl (664.8 MB)
 664.8/664.8 MB 1.3 MB/s eta 0:00:00
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl (211.5 MB)
 211.5/211.5 MB 4.7 MB/s eta 0:00:00
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl (56.3 MB)
 56.3/56.3 MB 14.0 MB/s eta 0:00:00
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl (127.9 MB)
 127.9/127.9 MB 7.9 MB/s eta 0:00:00
Downloading nvidia_cusparselt_cu12-0.6.2-py3-none-manylinux2014_x86_64.whl (207.5 MB)
 207.5/207.5 MB 11.1 MB/s eta 0:00:00

```

```

Downloading nvidia_cusparsel_cul2-12.5.1.170-py3-none-manylinux2014_x86_64.whl (207.5 MB)
207.5/207.5 MB 6.0 MB/s eta 0:00:00
Downloading nvidia_nccl_cul2-2.21.5-py3-none-manylinux2014_x86_64.whl (188.7 MB)
188.7/188.7 MB 7.7 MB/s eta 0:00:00
Downloading nvidia_nvjitlink_cul2-12.4.127-py3-none-manylinux2014_x86_64.whl (21.1 MB)
21.1/21.1 MB 104.7 MB/s eta 0:00:00
Installing collected packages: nvidia-nvjitlink-cul2, nvidia-nccl-cul2, nvidia-curand-cul2, nvidia-cufft-cul2, n
Attempting uninstall: nvidia-nvjitlink-cul2
Found existing installation: nvidia-nvjitlink-cul2 12.5.82
Uninstalling nvidia-nvjitlink-cul2-12.5.82:
Successfully uninstalled nvidia-nvjitlink-cul2-12.5.82
Attempting uninstall: nvidia-nccl-cul2
Found existing installation: nvidia-nccl-cul2 2.23.4
Uninstalling nvidia-nccl-cul2-2.23.4:
Successfully uninstalled nvidia-nccl-cul2-2.23.4
Attempting uninstall: nvidia-curand-cul2
Found existing installation: nvidia-curand-cul2 10.3.6.82
Uninstalling nvidia-curand-cul2-10.3.6.82:
Successfully uninstalled nvidia-curand-cul2-10.3.6.82
Attempting uninstall: nvidia-cufft-cul2
Found existing installation: nvidia-cufft-cul2 11.2.3.61
Uninstalling nvidia-cufft-cul2-11.2.3.61:
Successfully uninstalled nvidia-cufft-cul2-11.2.3.61
Attempting uninstall: nvidia-cuda-runtime-cul2
Found existing installation: nvidia-cuda-runtime-cul2 12.5.82
Uninstalling nvidia-cuda-runtime-cul2-12.5.82:
Successfully uninstalled nvidia-cuda-runtime-cul2-12.5.82
Attempting uninstall: nvidia-cuda-nvrtc-cul2
Found existing installation: nvidia-cuda-nvrtc-cul2 12.5.82
Uninstalling nvidia-cuda-nvrtc-cul2-12.5.82:
Successfully uninstalled nvidia-cuda-nvrtc-cul2-12.5.82
Attempting uninstall: nvidia-cuda-cupti-cul2
Found existing installation: nvidia-cuda-cupti-cul2 12.5.82
Uninstalling nvidia-cuda-cupti-cul2-12.5.82:
Successfully uninstalled nvidia-cuda-cupti-cul2-12.5.82
Attempting uninstall: nvidia-cublas-cul2
Found existing installation: nvidia-cublas-cul2 12.5.3.2
Uninstalling nvidia-cublas-cul2-12.5.3.2:
Successfully uninstalled nvidia-cublas-cul2-12.5.3.2
Attempting uninstall: nvidia-cusparse-cul2
Found existing installation: nvidia-cusparse-cul2 12.5.1.3
Uninstalling nvidia-cusparse-cul2-12.5.1.3:
Successfully uninstalled nvidia-cusparse-cul2-12.5.1.3
Attempting uninstall: nvidia-cudnn-cul2
Found existing installation: nvidia-cudnn-cul2 9.3.0.75
Uninstalling nvidia-cudnn-cul2-9.3.0.75:
Successfully uninstalled nvidia-cudnn-cul2-9.3.0.75
Attempting uninstall: nvidia-cusolver-cul2
Found existing installation: nvidia-cusolver-cul2 11.6.3.83
Uninstalling nvidia-cusolver-cul2-11.6.3.83:
Successfully uninstalled nvidia-cusolver-cul2-11.6.3.83
Successfully installed bitsandbytes-0.46.1 nvidia-cublas-cul2-12.4.5.8 nvidia-cuda-cupti-cul2-12.4.127 nvidia-cu

```

```

from huggingface_hub import notebook_login
notebook_login()

```



- Code below reference to: <https://huggingface.co/meta-llama/Llama-2-7b-chat-hf>

```

from huggingface_hub import login
login(new_session=False)

```

```
summary = "The CAD to CNY exchange rate is forecasted to decrease by 0.19% over the next 14 days, from 5.2318 to 5.1"
```

```
# Use a pipeline as a high-level helper
from transformers import pipeline
```

```
pipe = pipeline("text-generation", model="meta-llama/Llama-2-7b-chat-hf")
```

```
prompt = f"""You act a financial analyst, here is a exchange rate prediction made by a python project,
Please take a look at the result {summary}, and provide more insightful explanations.
"""
```

```
messages = [
    {"role": "user", "content": prompt},
]
pipe(messages)
```

 Loading checkpoint shards: 100% 2/2 [00:00<00:00, 1.91it/s]

Device set to use cuda:0

OutOfMemoryError Traceback (most recent call last)

/tmp/ipython-input-36575242.py in <cell line: 0>()

2 from transformers import pipeline

3

----> 4 pipe = pipeline("text-generation", model="meta-llama/Llama-2-7b-chat-hf")

5

6 prompt = f"""You act a financial analyst, here is a exchange rate prediction made by a python project,

⏮ 11 frames

/usr/local/lib/python3.11/dist-packages/torch/nn/modules/module.py in convert(t)

1327 memory_format=convert_to_format,

1328)

-> 1329 return t.to(

1330 device,

1331 dtype if t.is_floating_point() or t.is_complex() else None,

OutOfMemoryError: CUDA out of memory. Tried to allocate 86.00 MiB. GPU 0 has a total capacity of 14.74 GiB of which 82.12 MiB is free. Process 52947 has 14.66 GiB memory in use. Of the allocated memory 14.52 GiB is allocated by PyTorch, and 13.30 MiB is reserved by PyTorch but unallocated. If reserved but unallocated memory is large try setting PYTORCH_CUDA_ALLOC_CONF=expandable_segments:True to avoid fragmentation. See documentation for Memory Management (<https://pytorch.org/docs/stable/notes/cuda.html#environment-variables>)


```
from transformers import AutoTokenizer, AutoModelForCausalLM
import torch


model_name = "meta-llama/Llama-2-7b-chat-hf" # change to any model you want

tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(model_name)

messages = [
    {"role": "user", "content": "Who are you?"},
]

inputs = tokenizer.apply_chat_template(
    messages,
    add_generation_prompt=True,
    tokenize=True,
    return_dict=True,
    return_tensors="pt",
).to(model.device)

outputs = model.generate(**inputs, max_new_tokens=40)
print(tokenizer.decode(outputs[0][inputs["input_ids"].shape[-1]:]))
```



model.safetensors.index.json: 100%	26.8k/26.8k [00:00<00:00, 608kB/s]
Fetching 2 files: 100%	2/2 [07:47<00:00, 467.59s/it]
model-00001-of-00002.safetensors: 100%	9.98G/9.98G [07:47<00:00, 56.6MB/s]
model-00002-of-00002.safetensors: 100%	3.50G/3.50G [05:15<00:00, 6.42MB/s]
Loading checkpoint shards: 0%	0/2 [00:00<?, ?it/s]

the session crashed after using all available GPUs. So will try using Hugging Face Inference API call.

Using Inference API instead of local GPUs

- code reference for inference API: https://huggingface.co/meta-llama/Meta-Llama-3-8B-Instruct?inference_api=true&inference_provider=novita&language=python&client=openai

```
import os
os.environ['HF_TOKEN'] = ''

prompt = f"""You act a financial analyst, here is a exchange rate prediction made by a python project,
Please take a look at the result {summary}, and provide more insightful explanations.
"""
```

```
import os
from huggingface_hub import InferenceClient
```

```
client = InferenceClient(
    provider="novita",
    api_key=os.environ["HF_TOKEN"],
)
```

```
completion = client.chat.completions.create(
    model="meta-llama/Meta-Llama-3-8B-Instruct",
    messages=[
        {
            "role": "user",
            "content": prompt
        }
    ],
)
```

```
print(completion.choices[0].message)
```

```
↔ ChatCompletionOutputMessage(role='assistant', content="**Exchange Rate Forecast Analysis**\n\nThank you for providing the exchange rate prediction model output. As a financial analyst, I'll review the results and provide key findings and recommendations based on the forecasted data.
```

```
print(completion.choices[0].message['content'])
```

```
↔ **Exchange Rate Forecast Analysis**
```

Thank you for providing the exchange rate prediction model output. As a financial analyst, I'll review the results and provide key findings and recommendations based on the forecasted data.

****Key Findings:****

- **CAD to CNY Exchange Rate Decrease:**** The model predicts a decrease in the CAD to CNY exchange rate by 0.19% over the 14-day period.
- **Forecasted Exchange Rate:**** The forecasted exchange rate for CAD to CNY on 2025-08-14 is 5.2220, down from the current rate of 5.2310.
- **Timeframe:**** The prediction is made over a 14-day period, which is a relatively short-term forecast.

****Insights:****

- **Market Sentiment:**** The predicted decrease in the CAD to CNY exchange rate suggests a slightly bearish market sentiment for the CAD.
- **Risk Management:**** For investors and businesses with exposure to the CAD to CNY exchange rate, a 0.19% decrease could impact their financials.
- **Economic Indicators:**** To further understand the drivers behind this forecast, it's crucial to examine economic indicators for both Canada and China.
- **Monitoring and Re-evaluation:**** Given the relatively short-term nature of the forecast, it's essential to monitor the exchange rate closely and re-evaluate the forecast as more data becomes available.

****Recommendations:****

- **Continuously Monitor Exchange Rates:**** Keep a close eye on the CAD to CNY exchange rate and any changes in market sentiment.
- **Assess Economic Indicators:**** Examine economic indicators for both Canada and China to better understand the factors influencing the exchange rate.
- **Consider Hedging Strategies:**** If you have exposure to the CAD to CNY exchange rate, consider implementing hedging strategies to mitigate the risk of a decrease.
- **Re-evaluate the Forecast:**** Regularly re-evaluate the forecast as more data becomes available to refine the prediction.

By taking these steps, you can make more informed decisions and better manage the risks associated with exchange rate fluctuations.